

98-008 Homework 1: PPM Parsing

Cooper Pierce cppierce@andrew.cmu.edu
Jack Duvall jrduvall@andrew.cmu.edu

Fall 2022

Overview

The goal of assignment is to get you used to basic Rust constructs and syntax, filling out the core of a program that actually does something useful! We hope this will give you familiarity with Rust's development workflow and make you appreciate the language.

This assignment will have you write a basic parser for a simple image format known as PPM.

PPM Images

PPM is a very simple image format, consisting of just:

- A “magic number” to distinguish it from other files
- Width and height of the image
- The maximum intensity value of
- Pixels as packed RGB pixel values, read sequentially row-by-row from the top left of the image.

<http://ailab.eecs.wsu.edu/wise/P1/PPM.html> outlines the format; for simplicity, we'll summarize it again here. Note that we'll be using the binary format, where each pixel channel takes up exactly one byte.

PPM Format Specification

```
PPM ::= MagicHeader
      CommentsAndWhitespace Width
      CommentsAndWhitespace Height
      CommentsAndWhitespace Maxval
      “\n” Pixels
MagicHeader ::= “P6”
CommentsAndWhitespace ::= (“\t” | “\n” | “\x0C” | “\r” | “ ” | (“#” Comment “\n”))
                        CommentsAndWhitespace
Comment ::= ((Any character except for “\n”) Comment) | ε
Width ::= (“0” | “1” | “2” | “3” | “4” | “5” | “6” | “7” | “8” | “9”) Width | ε
Height ::= Width
Maxval ::= Maxval
Pixels ::= Exactly  $w * h * 3$  bytes, where  $w$  and  $h$  are the parsed width and height
```

How to read this table: groups being right next to each other means the text they match should be concatenated. The | operator means that either the left or right character can be accepted. Characters in “” means only that exact string is accepted. ϵ matches the empty string.

PPM Image Example

An example would probably help understand the above specification a bit better:

```
P6
# I'm a PPM file! You can tell by the magic header
3 3 # This first number is the width, and the second is the height
# This next number is the maximum intensity of each pixel
# Exporting as "raw" in GIMP always gives 255, which makes sense; this is the
# maximum value of a u8. After this maxval, we are only allowed a single newline
# before the pixels start. These pixels can be arbitrary bytes!
255
<<<<<<<<<aaaaaaaa~~~~~~
```

Copy-pasting this text into a file ending with .ppm and opening it in a compatible image-viewing program should give a 3x3 image with 3 horizontal gray stripes.